

# Introduction to Deep Learning

## Machine Learning

### Linear Regression

$$\hat{y}_i = \sum_{j=1}^d x_{ij} \theta_j$$

$d$  = input dimension

$x$  = input data

$\theta$  = weights

Loss function:

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

in matrix notation:

$$J(\theta) = (\mathbf{X}\theta - \mathbf{y})^T (\mathbf{X}\theta - \mathbf{y})$$

differentiate:

$$\frac{\partial J(\theta)}{\partial \theta} = 2\mathbf{X}^T \mathbf{X} \theta - 2\mathbf{X}^T \mathbf{y} = 0$$

$$\theta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Analytical solution for a convex problem.

### Logistic Regression

$$\hat{\mathbf{y}} = p(\mathbf{y} = 1 | \mathbf{X}, \theta) = \prod_{i=1}^n p(y_i = 1 | \mathbf{x}_i, \theta)$$

where

$$\hat{y}_i = \sigma(\mathbf{x}_i \theta)$$

$\sigma$  is the sigmoid/logistic function.

Loss function (binary cross entropy loss):

$$L(\hat{y}_i, y_i) = -[y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

No analytical solution. Iterative method needed (gradient descent).

## Loss Functions

Measure the goodness of the prediction.

Large loss indicates bad performance.

Choice of loss function depends on problem.

### Regression Loss

L1 Loss:

$$L(y, \hat{y}; \theta) = \frac{1}{n} \sum_i^n ||y_i - \hat{y}_i||_1$$

MSE Loss:

$$L(y, \hat{y}; \theta) = \frac{1}{n} \sum_i^n ||y_i - \hat{y}_i||_2^2$$

where

L1 norm (Manhattan distance):

$$||x||_1 = \sum |x_i|$$

- Promotes sparsity
- More robust to outliers

L2 norm (Euclidean distance):

$$||x||_2 = \sqrt{\sum x_i^2}$$

- Differentiable
- Penalizes large errors more heavily

Squared L2 norm:

$$||x||_2^2 = \sum x_i^2$$

- Differentiable

- Computationally efficient

## Classification Loss

Binary Cross Entropy loss (binary classification):

$$L(y, \hat{y}; \theta) = -\frac{1}{n} \sum_i^n [y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)]$$

$y$  is a probability output ( $[0, 1]$ ) of the first class. Therefore,  $1 - y$  is the probability of the second class. Works well with a sigmoid activation in the last layer.

Log penalizes bad predictions:  $\log 0 = -\infty$ ,  $\log 1 = 0$  (sort of inverts the prediction  $\hat{y}$ ).

Encourages the model to make confident predictions close to 0 or 1.

Cross Entropy loss (multi-class classification):

$$L(y, \hat{y}; \theta) = - \sum_{i=1}^n \sum_{k=1}^k (y_{ik} \cdot \log \hat{y}_{ik})$$

The second sum is a generalization of the binary cross entropy case for multiple classes.